

# AI Usage Audit

Eric Wang — Product Operations Internship, Instawork · May 1 – June 10, 2026 (first recorded usage May 5) · Claude Code

1.44B  
total tokens

\$1,625  
API-equivalent cost

6,781  
API calls

41  
interactive sessions



\$6,624  
saved by prompt caching

94.7%  
of tokens were cache reads

9,272  
tool invocations

34/37  
days active


Over 37 days, every working day was AI-assisted (34 active days, longest streak 16 consecutive days). The 6,781 model generations were chained by 9,272 tool executions — 2,521 shell commands, 1,676 file reads, 1,499 code edits, and 515 production Redshift queries. Peak single day: June 9 (\$257). All usage ran on a subscription, so marginal cost was \$0 — every dollar figure in this report is the API list-price equivalent.




Cost & token breakdown by model

| Model      | API calls | Input  | Output | Cache w (5m) | Cache w (1h) | Cache reads | Cost       |
|------------|-----------|--------|--------|--------------|--------------|-------------|------------|
| Opus 4.7   | 4,074     | 125.6K | 5.16M  | 5.83M        | 29.0M        | 830.3M      | \$871.30   |
| Opus 4.8   | 1,774     | 469.4K | 1.24M  | 5.30M        | 22.8M        | 357.2M      | \$472.71   |
| Fable 5    | 876       | 222.3K | 366.3K | 3.80M        | 1.78M        | 175.4M      | \$278.99   |
| Sonnet 4.6 | 8         | 17     | 2.99K  | 178.3K       | 92.1K        | 140.9K      | \$1.31     |
| Haiku 4.5  | 49        | 12.0K  | 9.74K  | 279.3K       | 0            | 1.94M       | \$0.60     |
| TOTAL      | 6,781     | 829.2K | 6.79M  | 15.4M        | 53.6M        | 1.36B       | \$1,624.91 |


Cost share by model




Token composition



**Monthly split:** May \$940 (5.45M output tokens) · June 1–10 \$685 (1.33M output). Average per active day: \$48. The model timeline tells the story of the internship: Opus 4.7 carried May, Opus 4.8 took over after its release, and Fable 5 arrived for the final stretch.




## Usage over time




## Activity patterns

Most active hour: **17:00–18:00** local time (1,117 API calls). Late-night work (10pm–6am): **2%** of all calls. Weekend share: **7%**.




**When the work happened (API calls, weekday x hour)**



## What the calls actually were

Each API call is one model generation. Interactive sessions chain generations through tool results; Agent-tool subagents and Workflow fan-outs (parallel multi-agent runs) execute their own chains underneath the main sessions.



## Most expensive sessions

Each row is one interactive session including every subagent and workflow agent it spawned. Titles are the session names from the transcript (or the opening request).

| Session                                                    | Date   | Calls | Agents | Output | Cost  |
|------------------------------------------------------------|--------|-------|--------|--------|-------|
| AI Agent Suite ReDesign                                    | May 27 | 2,252 | 47     | 696.5K | \$418 |
| captain dash revamp                                        | May 16 | 1,055 | 9      | 1.35M  | \$265 |
| capt dash inappchat                                        | May 26 | 350   | 0      | 665.0K | \$183 |
| capt promo automation replit                               | May 11 | 543   | 10     | 729.8K | \$152 |
| captains OH                                                | May 29 | 363   | 17     | 420.2K | \$102 |
| financial agent                                            | May 18 | 172   | 0      | 336.8K | \$76  |
| boson chatbot                                              | May 21 | 261   | 0      | 261.2K | \$72  |
| write this email out for me sending to uchicago robotics c | May 06 | 393   | 34     | 338.2K | \$56  |
| internship wrapup                                          | Jun 10 | 127   | 12     | 103.4K | \$55  |
| let's begin another manual captain search process for this | May 14 | 135   | 3      | 280.5K | \$33  |

## Week by week

| Week of | Active days | API calls | Output tokens | Cost     | Share |
|---------|-------------|-----------|---------------|----------|-------|
| May 04  | 5           | 342       | 364.6K        | \$45.96  | 2.8%  |
| May 11  | 7           | 1,278     | 1.69M         | \$280.24 | 17.2% |
| May 18  | 6           | 1,596     | 2.09M         | \$393.38 | 24.2% |
| May 25  | 6           | 1,166     | 1.30M         | \$220.25 | 13.6% |
| Jun 01  | 7           | 1,381     | 836.2K        | \$311.31 | 19.2% |
| Jun 08  | 3           | 1,018     | 497.2K        | \$373.78 | 23.0% |

## Methodology & pricing

**Source:** 232 Claude Code JSONL transcripts under `~/ .claude/projects/` — interactive sessions, Agent-tool subagents, and Workflow agents. Each assistant API message carries a usage block; multi-block messages repeat the same message id across lines, so messages were deduplicated by id (6,781 unique generations from 232 files). Timestamps converted from UTC to local time. Tool invocations counted from `tool_use` content blocks.

| Model          | Input \$/MTok | Output \$/MTok | Cache write 5m | Cache write 1h | Cache read |
|----------------|---------------|----------------|----------------|----------------|------------|
| Opus 4.7 / 4.8 | \$5.00        | \$25.00        | \$6.25         | \$10.00        | \$0.50     |
| Fable 5        | \$10.00       | \$50.00        | \$12.50        | \$20.00        | \$1.00     |
| Sonnet 4.6     | \$3.00        | \$15.00        | \$3.75         | \$6.00         | \$0.30     |
| Haiku 4.5      | \$1.00        | \$5.00         | \$1.25         | \$2.00         | \$0.10     |

## Caveats

- Figures are API list-price equivalents. Actual usage ran on a Claude subscription — marginal cost \$0.
- Excludes `claude.ai` web/mobile usage, cloud code-review runs (ultrareview), and any transcripts deleted before the audit.
- No transcripts exist before May 5; the May 1–4 contribution is zero recorded Claude Code usage.
- Cache write tokens are split 5-minute vs 1-hour TTL per the per-message usage breakdown and priced at 1.25x / 2.0x base input respectively; cache reads at 0.1x.
- Tool invocation counts come from transcript content blocks and may differ marginally from billed tool executions.